

Comparison of Regression Methods for Approximating Runge’s Function under Noisy Conditions

Live L. Storborg and Henrik Haug
*Department of Mathematics, University of Oslo,
PO Box 1053, Blindern 0316, Oslo, Norway*

AND

Simon S. Thommesen and Adam Falchenberg
*Institute of Theoretical Astrophysics, University of Oslo,
PO Box 1029, Blindern 0315, Oslo, Norway*

(Dated: October 6, 2025)

Approximating smooth functions with polynomial regression is prone to numerical instability known as Runge’s phenomenon, which worsens in the presence of noise. To address this, we compare Ordinary Least Squares (OLS), Ridge, and Lasso regression for approximating Runge’s function under noisy conditions using both analytical and gradient-based optimization, alongside bootstrap and cross-validation analyses. OLS exhibits clear overfitting beyond polynomial degree ten, while Ridge mitigates this through coefficient shrinkage, remaining stable up to degree thirty. Lasso provides the most robust performance by eliminating irrelevant features and maintaining low test MSE across all degrees for higher noise with standard deviation $\sigma \geq 0.2$. Resampling by bootstrapping identified optimal complexity around degree fourteen and 300 data points, while resampling by cross validation identified optimal complexity with respect to OLS to favor low sample sizes and higher polynomial degrees, except for when the noise is sufficiently low that OLS performs best. Overall, Lasso regression best balances bias and variance, effectively suppressing Runge’s phenomenon and demonstrating strong generalization across varying noise levels.

I. INTRODUCTION

Machine learning and statistical modelling face a fundamental challenge: balancing model complexity with predictive accuracy. Overly simple models fail to capture underlying patterns (high bias) while overly complex models overfit to training data noise (high variance). Regression analysis is key for understanding these relationships, with polynomial fitting representing a classic example.

The Runge function, defined as

$$f(x) = \frac{1}{1 + 25x^2} \quad \text{for } x \in [-1, 1], \quad (1)$$

provides an ideal test case for studying regression methods due to its association with Runge’s phenomenon. This phenomenon demonstrates that high-degree polynomial interpolation can experience severe oscillations near the boundaries, making it challenging to achieve stable approximations [1].

To simulate realistic experimental conditions, we add stochastic noise to the Runge function using a normal distribution:

$$y_i = f(x_i) + \epsilon_i, \quad \epsilon_i \sim \mathcal{N}(0, \sigma^2), \quad (2)$$

where ϵ_i represents random noise with zero mean and standard deviation σ . This noise addition mimics real-world measurement uncertainties and allows us to study how the regression methods handle noisy data.

In this project, we implement and compare three regression models: Ordinary Least Squares (OLS), Ridge regression, and Lasso regression. We explore how well these can approximate the Runge function through polynomial fitting. Our analysis includes a systematic investigation of hyperparameters such as regularization strength λ , learning rate η , polynomial degree, and bias-variance analysis using bootstrap resampling. We also implement various gradient descent optimization methods including momentum, AdaGrad, RMSprop, and ADAM to explore their effectiveness in finding optimal parameters.

In this article, section II describes the theoretical foundations of the regression methods and our implementation approach, including gradient descent algorithms and resampling techniques. Section III presents our experimental results, which are discussed in section IV where bias-variance analysis and systematic parameter optimization are studied to explore how well the models perform. Finally, Section V summarizes our key findings regarding optimal model selection for the Runge function and discusses how polynomial regression provides groundwork for more complex optimization problems.

II. METHOD & THEORY

We will assume we have some data \mathbf{y} described by the Runge function (1). To find an approximation $\tilde{\mathbf{y}}$ of the Runge function, we will apply polynomial regression. We

will analyze the three regression models; OLS, Ridge and Lasso. Each approach has its own cost function to determine the optimal parameters θ .

To study the performance of the different models we use the two evaluation metrics, the Mean squared error (MSE) and the R^2 -score function. The MSE measures the squared difference between the predicted and the true value:

$$MSE(\mathbf{y}, \tilde{\mathbf{y}}) = \frac{1}{n} \sum_{i=0}^{n-1} (y_i - \tilde{y}_i)^2. \quad (3)$$

The R^2 score, however, quantifies how well the model performs compared to a model that always predicts the mean value. Values range from 0 (model performs no better than the mean) to 1 (a perfect fit):

$$R^2(\mathbf{y}, \tilde{\mathbf{y}}) = 1 - \frac{\sum_{i=0}^{n-1} (y_i - \tilde{y}_i)^2}{\sum_{i=0}^{n-1} (y_i - \bar{y})^2}, \quad (4)$$

where \tilde{y}_i is the i -th predicted value, y_i is the corresponding true value, and \bar{y} is the mean value of \mathbf{y} , defined as

$$\bar{y} = \frac{1}{n} \sum_{i=0}^{n-1} y_i.$$

A. Creating the Design Matrix & Cost Function

For polynomial fitting, we seek to approximate the function using a polynomial of degree p :

$$\tilde{y}(x) = \theta_0 + \theta_1 x + \theta_2 x^2 + \dots + \theta_{p-1} x^{p-1} = \sum_{i=0}^{p-1} \theta_i x^i, \quad (5)$$

where $\theta = [\theta_0, \theta_1, \dots, \theta_{p-1}]^T$ are the polynomial coefficients to be determined.

Given n data points, the design matrix \mathbf{X} is constructed as the Vandermonde matrix

$$\mathbf{X} = \begin{bmatrix} 1 & x_1 & x_1^2 & \dots & x_1^{p-1} \\ 1 & x_2 & x_2^2 & \dots & x_2^{p-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_n & x_n^2 & \dots & x_n^{p-1} \end{bmatrix}, \quad (6)$$

where x_1, x_2, \dots, x_n are the input data points sampled uniformly from the interval $[-1, 1]$, and each row represents one data point with its corresponding polynomial features. In essence, the linear system $\mathbf{X}\theta = \mathbf{y}$ relates the design matrix (6) and coefficients to the target values contained in \mathbf{y} . When we cannot recreate the data perfectly using our model, the predicted values are given by the design matrix and the optimal parameters $\hat{\theta}$ found by the model, $\mathbf{X}\hat{\theta} = \tilde{\mathbf{y}}$.

The optimal parameters $\hat{\theta}$ are found by minimizing a cost function. The cost functions of the models are usually given by the MSE (3) (there will be some differences,

see section IID) which expresses the difference between the predicted values and the data. When discussing the role of the design matrix, we saw that the predicted values $\tilde{\mathbf{y}}$ can be expressed as a function of the parameters θ and the design matrix \mathbf{X} . Using this, as well as writing equation 3 in vector form, we can define the cost function

$$C(\theta) = \frac{1}{n} \{(\mathbf{y} - \mathbf{X}\theta)^T (\mathbf{y} - \mathbf{X}\theta)\}. \quad (7)$$

Here, \mathbf{y} is the data, while $\tilde{\mathbf{y}} = \mathbf{X}\theta$ is the predicted values from our model, given by the calculated parameters θ . Since the data and the design matrix are given, the cost function is a function of only the model parameters.

To find the parameters that minimize the cost function, we differentiate the cost function with respect to θ and set the derivative equal to zero:

$$\frac{\partial C(\theta)}{\partial \theta} = 0. \quad (8)$$

Solving this equation will result in a stationary solution. To determine which stationary solution we get, we can examine the second derivative of the cost function.

$$\frac{\partial}{\partial \theta} \frac{\partial C(\theta)}{\partial \theta^T} = H \quad (9)$$

Here, H is the Hessian matrix. The Hessian matrix, given in equation 10, is positive semi-definite. Because of this, we know that the cost function is convex. A convex function will not have maxima or saddle points, only minima.

$$H = \frac{2}{n} \mathbf{X}^T \mathbf{X} \quad (10)$$

High complexity models can lead to an ill-conditioned Hessian matrix. This is especially true in our case with the Runge function, since it is defined in $x \in [-1, 1]$. If we look at the Vandermonde matrix (6), we can see that for high degree polynomials the last columns can become very similar. For instance, in a 15th-degree polynomial, the columns corresponding to x^{14} and x^{15} take on values that are very close to zero for most x -values except near the boundaries. As a result, the columns become nearly linearly dependent.

While this is not a theoretical issue, it poses practical problems: even small numerical errors can cause large changes in the model. In other words, the design matrix has a large condition number, meaning the system is highly sensitive to round-off errors, which causes ill-conditioning. The Hessian matrix amplifies this issue since it effectively squares the condition number. Later, we will see that the methods require the inverse of the Hessian matrix. The ill-conditioned nature of the matrix can lead to unstable and inaccurate solutions due to rounding errors.

B. Data splitting

In machine learning, it is customary to split the data. The data is typically split into training and test data. The training data is used to create the model $\tilde{\mathbf{y}}$. This is the data that is applied in the regression methods. The essence of data splitting is to train our model on the training data and test its accuracy and generality on the test data. In doing this, we can measure the variation of bias and variance with respect to the complexity of the model. Usually, the training data is chosen to be approximately 2/3 to 4/5 of the data.

An advantage of splitting the data into training and test sets is that it allows us to detect overfitting. Overfitting happens when the model captures the noise of the training data, and deviates from the underlying pattern of our full dataset. The training data is used to determine the optimal parameters $\boldsymbol{\theta}$. These parameters are then applied to the test data to obtain an independent approximation of the underlying function. By comparing the MSE on the training and test sets, we can assess whether the model is overfitting for higher complexity models (like higher degree polynomials).

C. Scaling

In machine learning, different features often have vastly different units and numerical scales. Without proper scaling, algorithms frequently perform poorly since features with large numerical values, regardless of their physical significance, will dominate the optimization process and create numerical instabilities. For polynomial fitting, this issue becomes particularly pronounced with high-order features like x^{15} , which can lead to ill-conditioned matrices, numerical overflow, and poor convergence in gradient descent algorithms.

To address these challenges, we standardize the design matrix \mathbf{X} using z-score normalization and apply mean centering to the target variable \mathbf{y} . For the design matrix, each feature is transformed according to

$$\mathbf{X}_{norm} = \frac{\mathbf{X} - \boldsymbol{\mu}_X}{\boldsymbol{\sigma}_X}, \quad (11)$$

where $\boldsymbol{\mu}_X$ and $\boldsymbol{\sigma}_X$ are the column-wise means and standard deviations computed from the training data. To prevent division by zero when encountering near-constant features, any standard deviation below 10^{-8} is set to 1.0, effectively leaving such features unscaled.

The target variable is mean-centered by

$$\mathbf{y}_{centered} = \mathbf{y} - \bar{\mathbf{y}} \quad (12)$$

where $\bar{\mathbf{y}}$ is a vector where all the elements are the mean μ_y of the training targets. This centering ensures that the intercept term θ_0 represents the model's prediction when all standardized features equal zero.

A critical aspect of our scaling procedure is that all scaling parameters ($\boldsymbol{\mu}_X$, $\boldsymbol{\sigma}_X$, and μ_y) are computed exclusively from the training data and then applied to both training and test sets. This approach prevents data leakage, where information from the test set could inadvertently influence the model training process. Using test set statistics for scaling would allow the model to gain knowledge about the test distribution, leading to an overly optimistic and biased evaluation of model performance.

The reason we scale the data using z-score normalization is to avoid numerical instability problems when inverting the Hessian matrix. As discussed previously, we can encounter very small values in the design matrix because of the domain of the Runge function. This is countered by applying z-score normalization. The contribution of the standard deviation in equation 11 increases the column-values so that every column has standard deviation $\sigma = 1$. This is especially helpful for the higher degree values, as they are scaled to avoid numerical rounding errors.

D. Regression Analysis

1. Ordinary Least Squares Regression

For the OLS method, we seek to minimize the cost function

$$\|\mathbf{y} - \mathbf{X}\boldsymbol{\theta}\|_2^2 = \frac{1}{n} (\mathbf{y} - \mathbf{X}\boldsymbol{\theta})^T (\mathbf{y} - \mathbf{X}\boldsymbol{\theta}). \quad (13)$$

The gradient becomes

$$\nabla_{OLS}(\boldsymbol{\theta}) = \frac{2}{n} (\mathbf{X}^T \mathbf{X} \boldsymbol{\theta} - \mathbf{X}^T \mathbf{y}), \quad (14)$$

which set to zero, gives the following expression for the optimal parameters:

$$\hat{\boldsymbol{\theta}}_{OLS} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}. \quad (15)$$

2. Ridge Regression

Ridge regression adds a regularization penalty term to the OLS cost function:

$$C_{Ridge}(\boldsymbol{\theta}) = \frac{1}{n} \|\mathbf{y} - \mathbf{X}\boldsymbol{\theta}\|_2^2 + \lambda \|\boldsymbol{\theta}\|_2^2, \quad (16)$$

where $\lambda > 0$ is a hyperparameter that controls the strength of regularization. This modification ensures that the matrix $\mathbf{X}^T \mathbf{X} + n\lambda \mathbf{I}$ is always invertible, even when $\mathbf{X}^T \mathbf{X}$ is singular or near singular.

The optimal parameters are obtained by minimizing the Ridge cost function. Taking the gradient with respect to $\boldsymbol{\theta}$ yields:

$$\nabla_{Ridge}(\boldsymbol{\theta}) = \frac{2}{n} (\mathbf{X}^T \mathbf{X} \boldsymbol{\theta} - \mathbf{X}^T \mathbf{y}) + 2\lambda \boldsymbol{\theta}, \quad (17)$$

Setting this gradient to zero and rearranging:

$$\begin{aligned}\frac{2}{n}(\mathbf{X}^T \mathbf{X} \boldsymbol{\theta} - \mathbf{X}^T \mathbf{y}) + 2\lambda \boldsymbol{\theta} &= \mathbf{0} \\ \mathbf{X}^T \mathbf{X} \boldsymbol{\theta} + n\lambda \boldsymbol{\theta} &= \mathbf{X}^T \mathbf{y} \\ (\mathbf{X}^T \mathbf{X} + n\lambda \mathbf{I}) \boldsymbol{\theta} &= \mathbf{X}^T \mathbf{y}\end{aligned}$$

The optimal Ridge regression parameters are then given by:

$$\hat{\boldsymbol{\theta}}_{\text{Ridge}} = (\mathbf{X}^T \mathbf{X} + n\lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}, \quad (18)$$

where \mathbf{I} is the identity matrix. The addition of $n\lambda \mathbf{I}$ to the diagonal guarantees that all eigenvalues are at least $n\lambda > 0$, ensuring numerical stability and preventing overfitting by shrinking the parameter values.

3. Lasso Regression

Lasso regression changes the L_2 penalty in equation 16 with an L_1 penalty in the cost function, which yields:

$$C_{\text{Lasso}}(\boldsymbol{\theta}) = \frac{1}{n} \|\mathbf{y} - \mathbf{X}\boldsymbol{\theta}\|_2^2 + \lambda \|\boldsymbol{\theta}\|_1. \quad (19)$$

The L_1 penalty provides more shrinkage than the ℓ_2 that Ridge regression uses, shrinking variables exactly to zero. This gives Lasso regression the characteristic of variable selection, as many of the variables completely vanish during the shrinkage process. Unlike Ridge, the L_1 term is not differentiable at 0, meaning we do not have an analytical solution, and we have to approximate it numerically. This leads the use of a subgradient when minimizing the cost function:

$$\nabla_{\text{Lasso}}(\boldsymbol{\theta}) = \frac{2}{n} (\mathbf{X}^T \mathbf{X} \boldsymbol{\theta} - \mathbf{X}^T \mathbf{y}) + \lambda \text{sgn}(\boldsymbol{\theta}). \quad (20)$$

where $\text{sgn}(0) \in [-1, 1]$.

E. Gradient Descent

Since the Lasso cost function has no closed-form minimizer, we use Gradient Descent to approximate its minimizer. Consider minimizing an objective $F : \mathbb{R}^p \rightarrow \mathbb{R}$. A Gradient Descent method iteratively updates a parameter vector by moving in the direction of the steepest descent of F . Starting from an initial guess $\boldsymbol{\theta}_0$, we continue to compute new values

$$\boldsymbol{\theta}_{k+1} = \boldsymbol{\theta}_k - \gamma_k \nabla F(\boldsymbol{\theta}_k), \quad k \geq 0, \quad (21)$$

with $\gamma_k > 0$ being the learning rate (step length). This procedure produces a sequence of approximations, where the sequence ideally converges to a minimizer of F .

Convergence is declared when relative MSE change falls below 10^{-6} . Divergence detection uses multiple criteria to distinguish catastrophic failure from temporary oscillations: immediate termination if MSE becomes NaN/infinite or exceeds $10\times$ the initial value, or if MSE increases consecutively over five iterations while exceeding $1.5\times$ the initial value. These criteria ensure that only persistent instability (typically from excessively large learning rates) triggers early termination, while normal oscillatory behavior is allowed to converge.

The method correlating to the update rule in equation 21 is known as Steepest Descent, and is the simplest and most traditional variant of Gradient Descent. While the scheme is conceptually simple, the method has several limitations. In machine learning applications, we usually work with complex functions which are not necessarily convex (as in Lasso regression). The gradient descent scheme might get stuck in local minima “ravines” and fails to approximate the global minima. The method is also inherently limited by its sensitivity to the learning rate γ_k .

1. Momentum

To address some limitations of steepest descent, we can introduce momentum-based gradient descent. This is a method that uses the momentum as a memory of the direction we are moving in parameter space, and introduces a velocity vector

$$\mathbf{v}_t = \gamma \mathbf{v}_{t-1} + \nabla_{\boldsymbol{\theta}} E(\boldsymbol{\theta}_t),$$

where $0 \leq \gamma \leq 1$ is the momentum parameter and $E(\boldsymbol{\theta}_t)$ is the expected value of $\boldsymbol{\theta}_t$. The velocity vector is used for averaging recently encountered gradients. The update rule is then given by

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \eta_t \mathbf{v}_t, \quad (22)$$

where η_t is the learning rate.

This method helps to dampen oscillations and accelerates convergence along directions of the gradient. Instead of relying solely on the current gradient, the algorithm uses a combination of past gradients to build velocity, which can help it move more smoothly through narrow, curved valleys.

2. RMS prop

RMS prop keeps track of a moving average of the squared gradients. This allows the method to adaptively scale the learning rate for each parameter based on the historical magnitude of its gradient. We start by denoting $\mathbf{s}_t = \mathbb{E}[\mathbf{g}_t^2]$, and we continue to calculate

$$\mathbf{g}_t = \nabla_{\boldsymbol{\theta}} E(\boldsymbol{\theta}),$$

and

$$\mathbf{s}_t = \beta \mathbf{s}_{t-1} + (1 - \beta) \mathbf{g}_t^2.$$

The update rule for RMS prop is then given by

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \eta_t \frac{\mathbf{g}_t}{\sqrt{\mathbf{s}_t + \epsilon}}, \quad (23)$$

where β is a parameter that controls the averaging time of the second moment, η is the learning rate, and ϵ is a small regularization constant to prevent divergences. This method effectively reduces the learning rate for parameters with consistently large gradients, and increases it for those with smaller gradients. This helps stabilize training and often leads to faster convergence.

3. ADAM

Another momentum-based method for gradient descent is the ADAM optimizer. ADAM stands for Adaptive Moment Estimation and combines ideas from both momentum and RMSprop. It keeps track of an exponentially decaying average of past gradients (first moment) and an exponentially decaying average of past squared gradients (second moment). These are denoted as:

$$\mathbf{m}_t = \mathbb{E}[\mathbf{g}_t], \quad \mathbf{s}_t = \mathbb{E}[\mathbf{g}_t^2],$$

where $\mathbf{g}_t = \nabla_{\boldsymbol{\theta}} E(\boldsymbol{\theta}_t)$ is the gradient of the cost function with respect to the parameters at time step t .

The moment estimates are updated as:

$$\mathbf{m}_t = \beta_1 \mathbf{m}_{t-1} + (1 - \beta_1) \mathbf{g}_t, \quad \mathbf{s}_t = \beta_2 \mathbf{s}_{t-1} + (1 - \beta_2) \mathbf{g}_t^2,$$

where β_1 and β_2 are hyperparameters that control the rate of decay. Since these moment estimates will be initialized as zero vectors, they will have a bias towards zero. Therefore, ADAM performs a bias correction

$$\hat{\mathbf{m}}_t = \frac{\mathbf{m}_t}{1 - \beta_1^t}, \quad \hat{\mathbf{s}}_t = \frac{\mathbf{s}_t}{1 - \beta_2^t}.$$

The update rule for ADAM is then given by:

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \eta_t \frac{\hat{\mathbf{m}}_t}{\sqrt{\hat{\mathbf{s}}_t + \epsilon}}. \quad (24)$$

4. ADAGrad

Another adaptive method for gradient descent is the AdaGrad optimizer, short for Adaptive Gradient Algorithm. The main idea of AdaGrad is to adapt the learning rate for each parameter individually, based on the history of its gradients. Parameters that receive frequent updates will have their learning rate reduced, while parameters that are updated less often will retain a relatively larger

learning rate. This makes AdaGrad particularly effective for problems with sparse data.

Formally, let

$$\mathbf{g}_t = \nabla_{\boldsymbol{\theta}} E(\boldsymbol{\theta}_t)$$

be the gradient of the cost function with respect to the parameters at time step t . AdaGrad keeps track of the accumulated squared gradients for each parameter:

$$\mathbf{r}_t = \mathbf{r}_{t-1} + \mathbf{g}_t \circ \mathbf{g}_t,$$

where \circ denotes element-wise multiplication.

The update rule for AdaGrad is then given by:

$$\boldsymbol{\theta}_{t+1,i} = \boldsymbol{\theta}_{t,i} - \frac{\eta}{\sqrt{r_{t,i}} + \epsilon} g_{t,i}, \quad (25)$$

where η is the initial learning rate and ϵ is a small constant added for numerical stability.

While AdaGrad works well for sparse problems and eliminates the need to manually tune per-parameter learning rates, a drawback is that the accumulated squared gradients \mathbf{r}_t grow without bound, causing the effective learning rates to shrink monotonically. This can eventually make the algorithm stop progressing.

5. Stochastic Gradient Descent

Stochastic Gradient Descent (SGD) is one of the most central optimization algorithms in machine learning. Unlike standard gradient descent, which computes the gradient of the cost function using the entire dataset, SGD approximates the gradient by using only a single randomly chosen sample, or a small batch of samples, at each update step. This drastically reduces the computational cost per iteration and makes training feasible on large datasets. The randomness also injects noise into the updates, which can help the algorithm escape shallow local minima or saddle points.

The starting point is that the cost function can be expressed as a sum over n data points:

$$C(\boldsymbol{\theta}) = \frac{1}{n} \sum_{i=1}^n c_i(\mathbf{x}_i, \boldsymbol{\theta}),$$

so that the gradient can be written as

$$\nabla_{\boldsymbol{\theta}} C(\boldsymbol{\theta}) = \frac{1}{n} \sum_{i=1}^n \nabla_{\boldsymbol{\theta}} c_i(\mathbf{x}_i, \boldsymbol{\theta}).$$

Instead of summing over all n terms, SGD estimates this gradient using a mini-batch $B_k \subset \{1, \dots, n\}$ of size M . The update rule is then given by

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \eta_t \frac{1}{M} \sum_{i \in B_k} \nabla_{\boldsymbol{\theta}} c_i(\mathbf{x}_i, \boldsymbol{\theta}_t), \quad (26)$$

where η_t is the learning rate at iteration t , and B_k is sampled uniformly at random.

In our implementation, we sample mini-batches **without replacement** within each epoch. The dataset is shuffled once at the start of each epoch, then divided sequentially into batches. This ensures each sample is used exactly once per epoch, reducing variance compared to sampling with replacement.

Convergence and divergence detection follow the same criteria as standard gradient descent (Section II E), but MSE is evaluated after each epoch rather than after each iteration. While SGD greatly improves efficiency, the noisy updates may lead to unstable convergence.

F. Resampling

Resampling methods are an important aspect of machine learning. By resampling, we can retrain our model on shuffled data and gain access to new estimations that can better our model performance and provide more reliable measures of its ability to generalize.

1. Bootstrap

Bootstrapping is a statistical resampling method that provides deeper insight into the variability and reliability of data-driven estimates. The idea of bootstrapping simply is to create new bootstrap datasets by sampling with replacement from our original dataset. These bootstrap samples are then used to recalculate the statistic of interest (such as the mean, median, or regression coefficients). By repeating this procedure many times, we can create a distribution of the calculations and compare it against the statistic of interest.

2. Cross-validation

Cross validation is another resampling technique we use to estimate how good our model fits the data. The idea of cross validation is to split our dataset into a number of folds. In each round, one fold will be used for test data, and the remaining folds will be used for training data. We will record how accurate our model's prediction is against the test data. The model is trained and evaluated on these splits repeatedly, rotating the role of the test fold each time.

By averaging performance over all folds, cross validation gives us a more reliable estimate of how well our model generalizes to new data. By running a k-fold cross validation technique on multiple machine learning methods, we will have obtained a good indication of which model to select, based on their ability to generalize well to unseen data.

3. Bias-Variance Tradeoff

In model selection, we are interested in knowing which model has a better fit to unseen data. Bias-Variance tradeoff is a central concept in how we would select the best model and why it has the best fit. Ideally, we would like a model with both low bias and low variance. In practice, however, this is rarely achievable. On one hand, simpler models are prone to have a high bias, and on the other hand, more complex models are prone to have a higher variance. The key to model selection is to find a good balance between bias and variance, as you would want to have a model that is complex enough to fit your data, but also simple enough to not overfit it.

4. Bias Variance Decomposition

We assume that the expectation values are taken over different realizations of the training data (e.g., through bootstrap sampling). For a fixed x , we study how the model prediction \tilde{y} varies when trained on different datasets. Starting from the mean squared error, the expected prediction error can be decomposed into bias, variance, and an irreducible noise term, where the noise is governed by a Gaussian distribution.

$$E[(y - \tilde{y})^2] = \text{Var}[\tilde{y}] + \text{Bias}_f^2[\tilde{y}] + \sigma^2. \quad (27)$$

Here, $\text{Var}[\tilde{y}]$ quantifies how sensitive the model is to variations in the training data, $\text{Bias}_f^2[\tilde{y}]$ measures the systematic error between the expected model prediction and the true function $f(x)$, and σ^2 represents the irreducible noise inherent in the data. A full step-by-step derivation of this result, including all intermediate algebraic expansions and expectation operations, is provided in Appendix A 1

G. Use of AI Tools

To enhance the quality of this scientific report, we utilized several AI tools throughout the writing and development process. Claude and ChatGPT were used to assist with formulating precise scientific language and improving sentence structure to meet academic writing standards. For code development, we used Claude, ChatGPT, and GitHub Copilot to support debugging processes, optimize code structure, and ensure consistency across our implementations. Additionally, these tools helped generate figures of good quality that maintain visual consistency throughout the document.

H. Tools

For code collaboration, we have used GitHub, and our code can be found by following this link [2]. Our code is

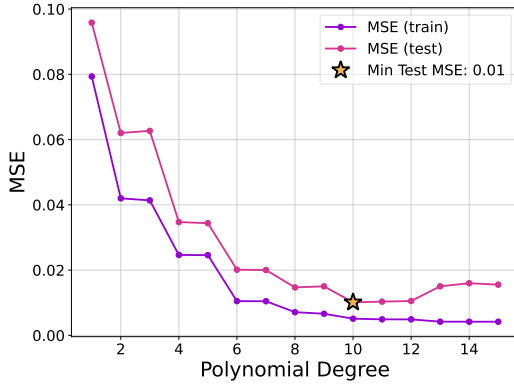


Figure 1. MSE of OLS training and test data for polynomial degree $p = 1, \dots, 15$ using 50 datapoints.

written in Python, with libraries: matplotlib [3], numpy [4], scikit-learn [5], seaborn [6] and pandas [7].

III. RESULTS

In this section, we will present how the regression models OLS, Ridge and Lasso, as well as the gradient descent methods, perform when approximating the Runge function with noise. We observe the results for different ranges of polynomial degrees and sample size. The noise added to the Runge function follows a normal distribution with mean $\mu = 0$ and standard deviation $\sigma = 0.1$. The data has been split into 3/4 training data and 1/4 test data. The scaling is done using z-score.

A. Ordinary Least Squares

We have implemented the OLS regression method for approximating the Runge function for polynomials up to degree 15, with data consisting of 34 to 1000 datapoints. In figure 1 we have plotted the MSE as a function of polynomial degree for 50 datapoints. The red curve shows the calculated MSE for the test data, while the blue curve shows the MSE for the training data. In figure 2 we have plotted the R^2 score for the models. The colors represent the same as in figure 1. Figure 3 shows the test MSE of the model as a function of polynomial degree and sample size. The heatmap shows the value of the MSE, where yellow represents high MSE values and blue represents low MSE values. The spread of values of the optimal parameters $\hat{\theta}$ are shown in figure 4 as a function of polynomial degree.

B. Ridge

Alongside the OLS method, we implemented Ridge regression to approximate the Runge function. Since Ridge

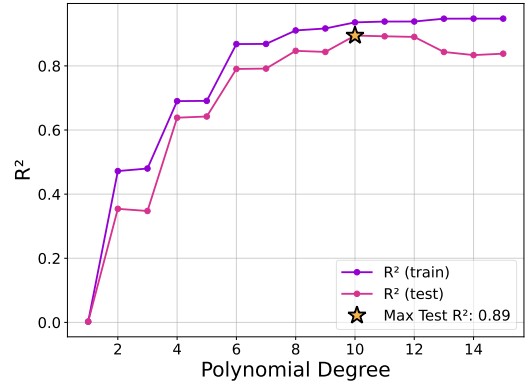


Figure 2. R^2 score of OLS training and test data for polynomial degree $p = 1, \dots, 15$ using 50 datapoints.

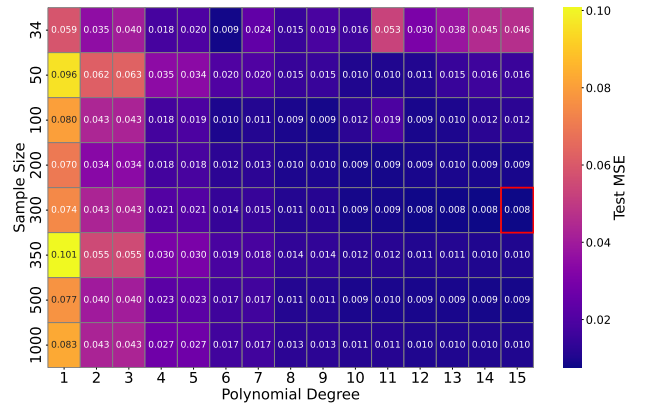


Figure 3. Test MSE of the OLS model as a function of polynomial degree and sample size. The red outline contains the lowest MSE.

regression introduces a penalty term that shrinks the regression coefficients, we extended the analysis to higher polynomial degrees to better illustrate its stability compared to OLS. In figure 5, we present the MSE for polynomial degrees up to 35, using 50 data points and a regularization parameter of $\lambda = 0.01$. This highlights when overfitting begins to appear and demonstrates how Ridge mitigates this effect. Similarly, figure 6 shows the corresponding R^2 scores for the same setup. Figure 7 illustrates the evolution of the regression coefficients θ as a function of the polynomial degree in the Ridge method. Lastly, figure 8 presents the relationship between polynomial degree, from 1 to 15, regularization parameter λ from $\log_{10}(-5)$ to $\log_{10}(2)$, sample size from 0 to 800 and shows the MSE for all the points in a 3D graph.

C. Gradient Descent

We evaluated five gradient descent optimizers (plain GD, Momentum, AdaGrad, RMSprop, and Adam) on OLS and Ridge regression ($\lambda = 0.01$). For each opti-




Figure 4. Spread of the values of the optimal parameters $\hat{\theta}$ of the OLS model as a function for polynomial degree $p = 1, \dots, 15$.

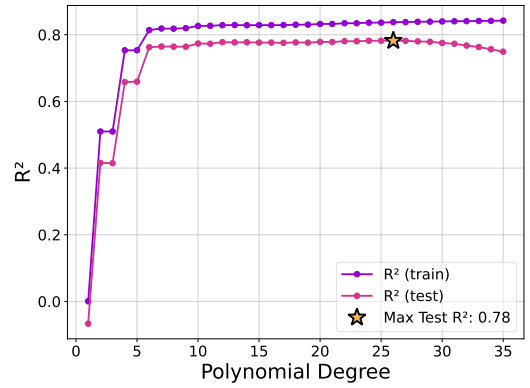


Figure 6. R^2 score of Ridge training and test data for polynomial degree $p = 1, \dots, 35$ using 50 datapoints with $\lambda = 0.01$.

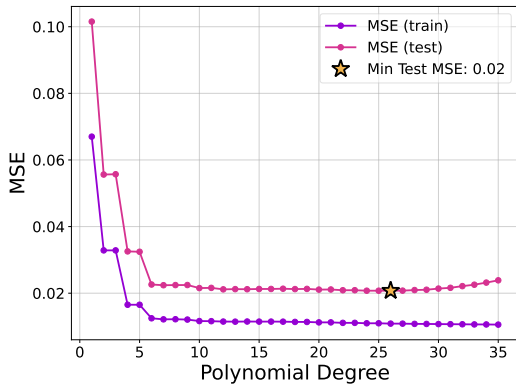


Figure 5. MSE of Ridge training and test data for polynomial degree $p = 1, \dots, 35$ using 50 datapoints with $\lambda = 0.01$.




Figure 7. Spread of the values of the optimal parameters $\hat{\theta}$ of the Ridge model as a function for polynomial degree $p = 1, \dots, 15$.

mizer, we tested multiple learning rates to identify values achieving convergence within 1500 iterations.

Figure 9 illustrates the learning rate sensitivity of steepest descent. Both OLS and Ridge diverged at $\eta = 0.37$ but converged at $\eta = 0.365$, demonstrating the critical importance of proper learning rate selection. Ridge consistently required more iterations than OLS (88 vs 93 iterations at $\eta = 0.365$), reflecting the additional computational cost of regularization.

Table I presents learning rate sensitivity for OLS across all optimizers. Momentum achieved fastest convergence (71 iterations), while RMSprop required the smallest learning rates and showed highly variable iteration counts. All converged methods achieved similar final MSE values (≈ 0.026).

Table II shows corresponding Ridge regression results. Compared to OLS, Ridge required more iterations across optimizers and exhibited increased sensitivity—notably, RMSprop failed to converge for $\eta \geq 0.01$. Final MSE values (≈ 0.027) were consistently higher than OLS due to the regularization penalty.

Table III presents optimal learning rates for Lasso regression. Momentum converged fastest (70 iterations),




Figure 8. Sample sizes 50 to 750, degrees 1 – 15, λ from $10^{-5} - 10^2$. Showing optimal MSE at degree 9, sample size 100 and $\lambda = 10^{-2}$.




Figure 9. Gradient descent convergence for OLS (solid) and Ridge (dashed, $\lambda = 0.01$) on the Runge function ($N = 500$, $d = 5$) with $\eta \in \{0.1, 0.2, 0.365, 0.37\}$. Vertical lines indicate convergence. Log-log scale.

while Lasso generally tolerated higher learning rates than Ridge (e.g., plain GD: $\eta = 0.5$ vs $\eta = 0.365$). Final MSE values ranged from 0.026 to 0.027.

Table I. OLS regression learning rate sensitivity ($N = 500$, $d = 5$) for five gradient-based optimizers.

Method	η	Iterations	Final MSE	Converged
GD	0.370	108	–	Diverged
GD	0.365	93	0.026051	Yes
GD	0.200	245	0.025893	Yes
GD	0.100	450	0.025901	Yes
Momentum	0.400	115	0.025886	Yes
Momentum	0.300	93	0.025887	Yes
Momentum	0.100	71	0.025890	Yes
Momentum	0.050	95	0.025890	Yes
AdaGrad	0.300	163	0.025890	Yes
AdaGrad	0.250	161	0.025890	Yes
AdaGrad	0.150	170	0.025890	Yes
AdaGrad	0.100	232	0.025893	Yes
RMSProp	0.020	136	0.027134	Yes
RMSProp	0.010	112	0.026289	Yes
RMSProp	0.005	185	0.025995	Yes
RMSProp	0.001	757	0.025892	Yes
Adam	0.200	99	0.025889	Yes
Adam	0.100	92	0.025896	Yes
Adam	0.050	86	0.025900	Yes
Adam	0.020	181	0.025888	Yes

D. Lasso

We first evaluated five gradient descent optimizers for Lasso regression with $\lambda = 0.01$ on a polynomial of degree $d = 5$, testing multiple learning rates to achieve convergence within 1500 iterations. Table III presents these results, demonstrating that gradient descent methods can successfully solve the non-differentiable Lasso optimization problem. Momentum and Adam achieved fastest convergence (70-84 iterations).

Table II. Ridge regression learning rate sensitivity ($N = 500$, $d = 5$, $\lambda = 0.01$) for five gradient-based optimizers.

Method	η	Iterations	Final MSE	Converged
GD	0.370	102	–	Diverged
GD	0.365	88	0.027516	Yes
GD	0.200	315	0.027038	Yes
GD	0.100	566	0.027051	Yes
Momentum	0.220	110	0.026993	Yes
Momentum	0.200	66	0.026804	Yes
Momentum	0.150	131	0.027018	Yes
Momentum	0.050	122	0.027043	Yes
AdaGrad	0.225	208	0.027033	Yes
AdaGrad	0.220	207	0.027033	Yes
AdaGrad	0.215	207	0.027033	Yes
AdaGrad	0.150	218	0.027034	Yes
RMSProp	0.010	1500	0.027355	No
RMSProp	0.005	162	0.027365	Yes
RMSProp	0.002	346	0.027159	Yes
RMSProp	0.001	648	0.027090	Yes
Adam	0.250	112	0.027059	Yes
Adam	0.200	98	0.027021	Yes
Adam	0.100	135	0.027034	Yes
Adam	0.050	119	0.027057	Yes

For the comparison across polynomial degrees shown in Figure 10, we used analytical solutions for OLS and Ridge, and scikit-learn’s coordinate descent implementation for Lasso to ensure robust convergence across all complexity levels.

Table III. Optimal learning rates and convergence behavior for Lasso regression ($N = 500$, $d = 5$, $\lambda = 0.01$) using five gradient-based optimizers. All methods successfully converged despite the non-differentiable L_1 penalty.

Method	η	Iterations	Final MSE
GD	0.500	171	0.027131
Momentum	0.050	70	0.025916
AdaGrad	0.200	231	0.026420
RMSprop	0.010	123	0.026228
Adam	0.100	84	0.025908

E. Stochastic Gradient Descent

In table IV we can see differences in calculation time, epochs/iterations, and MSE of Full- and Mini-Batch SGD for all the methods. We can see that the momentum method is both the fastest and the method that requires the least amount of epochs/iterations for both Full- and Mini-Batch. Outside this, there is more variation between the methods and batch type, especially for RMSprop and Adam. The final MSE of every method goes toward a similar value $MSE \simeq 0.027$.

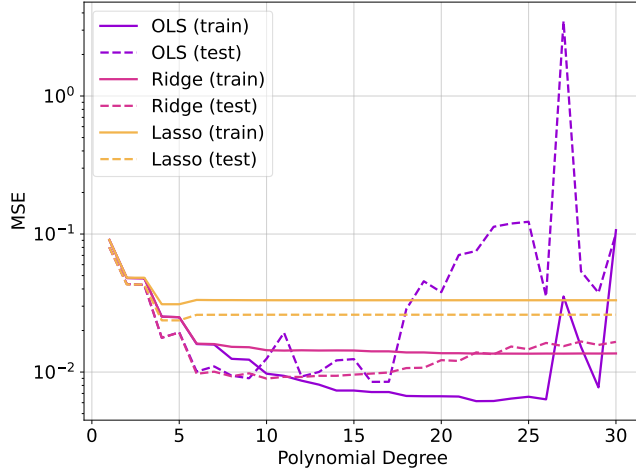


Figure 10. Comparison of training and test MSE for OLS, Ridge, and Lasso across polynomial degrees $p = 1, \dots, 30$ using 100 datapoints. OLS and Ridge use analytical solutions; Lasso uses scikit-learn’s coordinate descent.

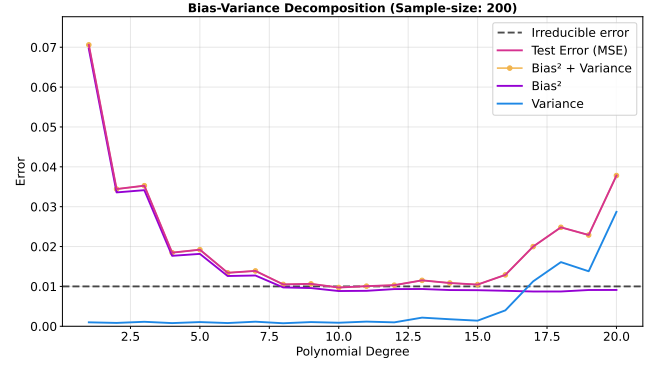


Figure 11. Bias-variance decomposition for OLS regression on Runge’s function with $n = 200$ data points and $\sigma = 0.1$. The blue curve shows the averaged variance, purple shows the averaged bias-squared, and yellow shows the sum of variance and bias-squared (the theoretical expected error/MSE). The pink curve is the averaged test error/MSE. The black dashed line indicates the irreducible error ($\sigma^2 = 0.01$).

Table IV. Convergence Comparison: Full-Batch vs Mini-Batch SGD for OLS ($N = 1,000,000$) using 40 epochs and batch size 256.

Method	Batch Type	Time (s)	Epochs/Iters	Final MSE
GD	Full	0.44	129	0.026848
Momentum	Full	0.25	71	0.026827
AdaGrad	Full	0.56	163	0.026829
RMSprop	Full	0.38	111	0.027161
Adam	Full	0.32	92	0.026835
SGD	Mini	0.40	11	0.026825
Momentum	Mini	0.27	7	0.026827
AdaGrad	Mini	0.48	12	0.026825
RMSprop	Mini	1.59	38	0.026826
Adam	Mini	0.79	17	0.026825

F. Bootstrapping

We used the bootstrapping resampling technique described in the theory to analyze the bias-variance decomposition for models with varying complexity; varying sample sizes n and polynomial degrees. In our analysis we used 2000 bootstrap samples. In 11 we present the bias variance decomposition for OLS using bootstrap, while in 12 we present a bootstrap bias-variance analysis for different sample sizes and model complexities. Finally, in 13, we represent different model predictions for different combinations of sample sizes and polynomial degree.

G. Cross-Validation

We performed 5-fold cross-validation resampling to compare the estimated model generalization error with the bootstrap estimated error. We then leveraged the CV method to compare the effect of regularization in Ridge and Lasso against OLS. The results are presented in Figures 14 and 15.




Figure 14. Bootstrap vs. cross-validation MSE comparison for OLS ($n = 200$, $\sigma = 0.1$). Pink: bootstrap test error (2000 resamples), purple: 5-fold CV validation error with ± 1 standard deviation (shaded).




Figure 12. Bootstrap bias-variance analysis across model complexities. Left panel shows MSE, middle panel shows bias^2 , and the right panel shows variance as functions of sample size n (y-axis) and polynomial degree (x-axis). The red boxes mark the model complexity ($n = 300$, degree = 14) that resulted in the minimum MSE. The coloring of the cells represents the magnitude of the MSE, bias^2 , and variance for each combination of model complexity.

IV. DISCUSSION

A. Ordinary Least Squares

We will discuss figures 1 and 2. We can see that as we increase the complexity, the MSE decreases for the training data. This is expected as the higher complexity models will be good at approximating the training data. The MSE of the test data also follows this trend for the lower degree polynomials. The test MSE is generally higher than the training MSE. This is consistent with the idea that the test data provides an independent measure of the model.

For the lowest polynomial degrees, we can see that the MSE is quite high, which corresponds to underfitting (not capturing the detail of the data). We can also see that the test MSE increases again after polynomial degree ten. This is overfitting. As described in section II B, this represents the fact the model, trained on the training data, is getting too detailed. This entails that the model deviates from the underlying pattern because the model starts memorizing the noise. We can see the same trend in figure 2 for the R^2 score of the model for 50 datapoints. The test data's R^2 score decreases after degree 10, while the test data's R^2 score goes towards one.

In figure 3 we can see that the MSE is not necessarily better for higher sample size. This is caused by the randomness of the noisy data. See section IV G for more clear information. In essence, what happens is that we hit a plateau where the MSE will be around the variance $\sigma^2 = 0.01$ introduced in the Gaussian noise. The reason the MSE is lower than this at some points comes from that the random data gives us a “lucky” dataset for these parameters.

B. Ridge Performance Analysis

Alongside the OLS regression method, we have also implemented Ridge regression to highlight the stabilizing effect of regularization. As shown in Figure 7, the penalty term introduced by Ridge shrinks the regression coefficients compared to the OLS model as seen in figure 4. This shrinkage reduces sensitivity to oscillations, particularly near the boundaries, and effectively mitigates Runge's phenomenon observed in the OLS case. By discouraging large coefficient values, Ridge regression produces a smoother and more stable approximation.

From a performance perspective, this regularization introduces a trade-off. Although Ridge slightly reduces the model's flexibility, it often improves generalization to unseen data. Compared to OLS, Ridge yields lower variance in prediction accuracy and maintains stability across a wider range of polynomial degrees.

This is evident in the comparison between Figures 1 and 5. The OLS model begins to show instability already around polynomial degree 12, whereas Ridge remains stable up to around degree 30—a clear advantage of the penalty term. Similarly, in Figures 2 and 6, the R^2 -score of Ridge demonstrates greater stability at higher polynomial degrees, reinforcing the benefits of regularization.

C. Gradient Descent Performance Analysis

To validate our gradient descent implementation and understand its behavior, we examined convergence patterns for OLS and Ridge regression across different learning rates: $\eta \in \{0.37, 0.365, 0.2, 0.05\}$. With properly tuned learning rates, gradient descent successfully reproduces the analytical solutions from Parts a) and b) within numerical precision. For instance, at $\eta = 0.365$, OLS converged to $\text{MSE} = 0.026051$ in 93 iterations, matching the analytical result (slightly higher) presented in figure

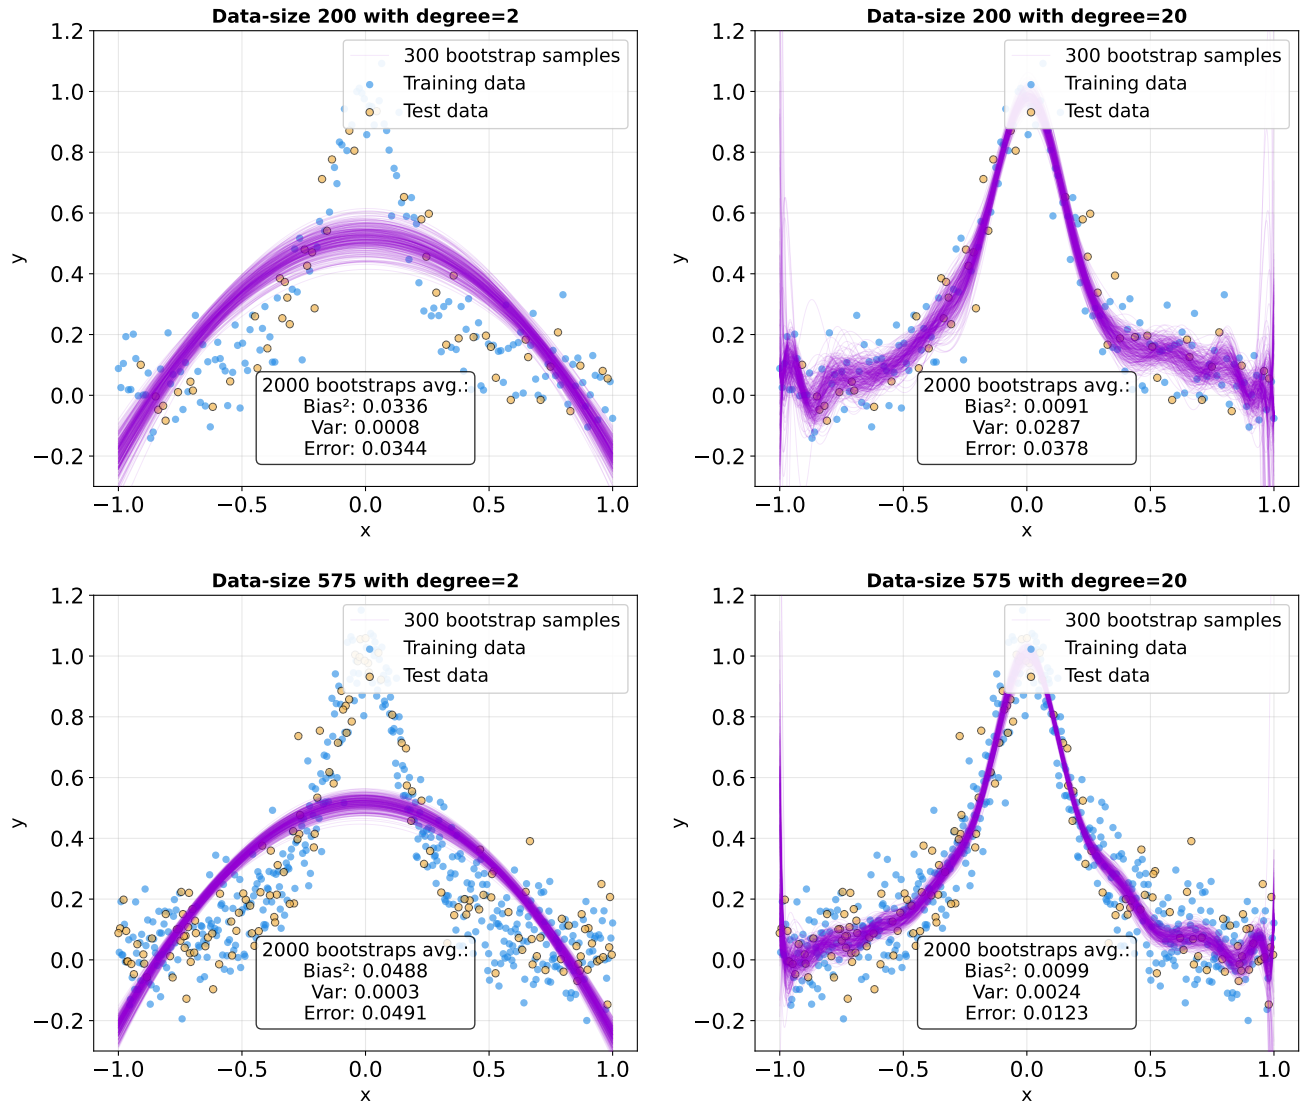


Figure 13. Model variance for individual bootstrap samples. Each subplot shows 300 bootstrap model predictions (purple, semi-transparent) for different combinations of sample sizes and polynomial degree. The bootstrap model used a train-test split of 25% test samples (yellow dots) and 75% training samples (blue dots).

3 for $N = 500$ and degree 5. While the analytical approach provides instantaneous solutions through direct matrix inversion, gradient descent requires iterative optimization but offers crucial flexibility for problems where closed-form solutions don't exist, as presented in IID 3, and will be discussed further in IV E.

Figure 9 demonstrates the critical importance of learning rate selection. At $\eta = 0.37$, both methods diverge, but reducing η by merely 0.005 to 0.365 achieves stable convergence. Ridge converged in 88 iterations (MSE = 0.027516) and OLS converged in 93 iterations (MSE = 0.026051). This sensitivity underscores that gradient descent requires careful tuning: learning rates that are too large cause divergence, while excessively small values

lead to prohibitively slow convergence.

Comparing OLS and Ridge, we observe that Ridge consistently requires slightly more iterations to converge. This is expected behavior stemming from the regularization term $\lambda \|\boldsymbol{\theta}\|_2^2$ in equation 16, which creates a "restoring force" that simultaneously fits the data while shrinking parameters (as discussed in IV B). However, this iteration penalty is offset by Ridge's improved numerical stability: the regularized matrix $\mathbf{X}^T \mathbf{X} + n\lambda \mathbf{I}$ (equation 18) is better conditioned than $\mathbf{X}^T \mathbf{X}$ (equation 15), enabling Ridge to use larger learning rates than OLS. This advantage should become more pronounced as model complexity increases and OLS's conditioning deteriorates.

Testing this hypothesis, we repeated the analysis at

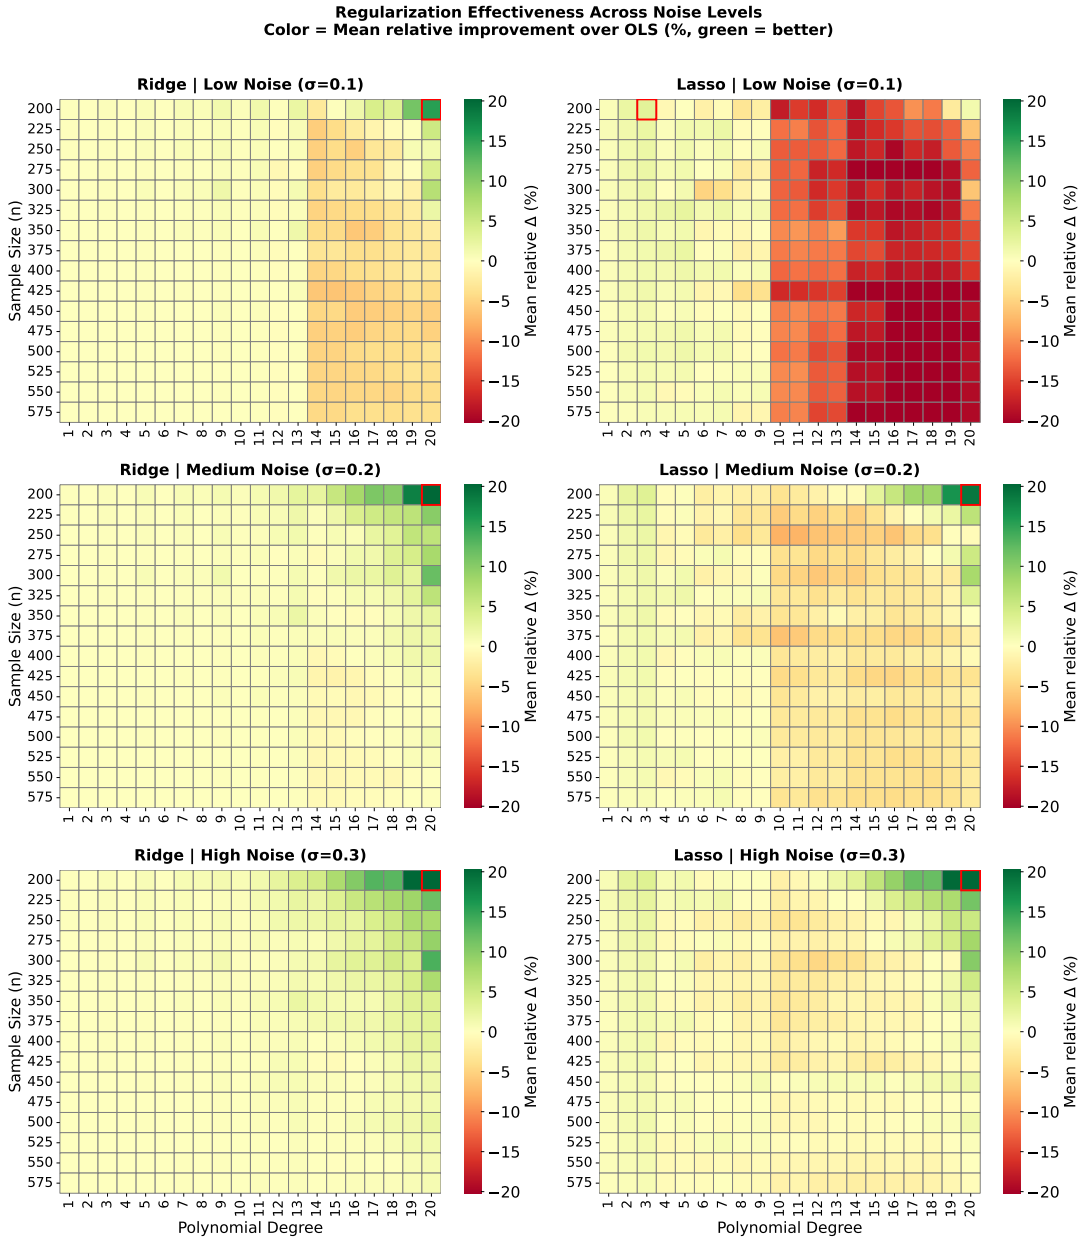


Figure 15. Relative improvement of regularization over OLS across noise levels. Each cell shows the mean relative improvement $\Delta_{\text{rel}} = (\text{MSE}_{\text{OLS}} - \text{MSE}_{\text{model}}) / \text{MSE}_{\text{OLS}}$ (in percent) averaged over 10 random seeds. Green indicates regularization improves over OLS, red indicates degradation. Left column shows Ridge (best λ), right column shows Lasso (best λ). Red boxes mark the maximum improvement for each configuration.

higher polynomial degrees (10 and 15). As model complexity increased, OLS gradient descent became increasingly unstable, requiring progressively smaller learning rates and more iterations to converge. Ridge regression, by contrast, maintained efficient convergence when λ was scaled appropriately with model complexity—using learning rates an order of magnitude larger than OLS while achieving convergence in comparable or fewer iterations. This confirms that Ridge’s regularization provides both statistical benefits (preventing overfitting) and computational advantages (enabling efficient gradient-based

optimization for ill-conditioned problems).

D. Gradient Descent Optimizers

Having established that plain gradient descent requires careful learning rate tuning, we now evaluate four adaptive optimization methods: Momentum, AdaGrad, RM-Sprop, and Adam. Tables I and II present comprehensive learning rate sweeps for each optimizer on both OLS and Ridge regression.

Momentum achieved the fastest convergence (71 iterations for OLS, 66 for Ridge), validating its theoretical advantage of velocity accumulation in consistent gradient directions. Adam demonstrated the second-best performance (86-99 iterations for OLS, 98-135 for Ridge) with notable stability across learning rates. AdaGrad was consistently slowest (161-232 iterations for OLS, 207-218 for Ridge) due to its monotonically decreasing effective learning rate, though it showed remarkable consistency. RMSprop exhibited highly variable behavior, with iteration counts ranging from 112 to 757 for OLS and failing to converge entirely at $\eta \geq 0.01$ for Ridge.

Examining learning rate robustness, Momentum tolerated the widest range ($\eta \in [0.05, 0.4]$) for OLS), while RMSprop showed extreme sensitivity, particularly for Ridge regression. Adam and AdaGrad demonstrated moderate robustness, though AdaGrad required more careful initial tuning for Ridge. Ridge consistently required more iterations than OLS across all optimizers (e.g., 88 vs 93 for plain GD at $\eta = 0.365$), reflecting the computational cost of the ℓ_2 penalty term. More critically, Ridge amplified sensitivity to learning rate selection—RMSprop’s convergence failure at moderate learning rates demonstrates this increased difficulty. All converged methods achieved similar final MSE values: approximately 0.026 for OLS and 0.027 for Ridge, with the higher Ridge MSE reflecting the expected bias-variance tradeoff from regularization.

These results demonstrate that Momentum offers the best combination of speed and robustness for the Runge function, while Adam provides reliable performance with minimal tuning. RMSprop’s poor performance on Ridge regression suggests caution when applying it to regularized problems.

In the case of OLS and Ridge regression, we have deterministic smooth cost functions (equations 13 and 16) that are differentiable everywhere. Adaptive optimizers are primarily designed to handle noisy gradients from stochastic mini-batch sampling, where gradient estimates vary between iterations. Our full-batch results show that these adaptive methods provide only modest improvements over plain gradient descent—the main advantage being increased robustness to learning rate selection rather than dramatic speedup in convergence. With proper tuning of η , plain gradient descent proves equally effective for these smooth optimization problems, as evidenced by the competitive performance of plain GD (93 iterations at $\eta = 0.365$) compared to Momentum (71 iterations).

E. Performance analysis of Lasso

Table III demonstrates that gradient descent methods can optimize Lasso effectively despite its non-differentiable penalty. Momentum-based methods (Momentum, Adam) converged fastest, while all optimizers achieved similar final MSE values (0.0259 – 0.0271), con-

firmed robust convergence to the same solution. Interestingly, the optimal learning rates for Lasso are notably higher than those for Ridge (e.g., plain GD: $\eta = 0.5$ vs. $\eta = 365$). This is likely due to Lasso’s constant-magnitude penalty gradient $\lambda \cdot \text{sgn}(\theta)$, whereas Ridge’s penalty gradient ($2\lambda\theta$) grows unboundedly with parameter magnitude, making it more sensitive to large learning rates.

Figure 10 reveals Lasso’s key advantage: stability at high polynomial degrees. While OLS and Ridge severely overfit with test MSE exploding while training MSE remains low. Lasso maintains low MSE for both training and test data throughout. Although Lasso’s training MSE is slightly higher than Ridge’s, it generalizes far better to unseen data.

This stability stems from Lasso’s feature selection capability. By setting irrelevant coefficients to exactly zero rather than shrinking them as Ridge does, Lasso eliminates high-degree terms that overfit to noise. Ridge keeps all terms active with small weights, allowing unnecessary high-degree features to degrade performance at high complexity. For the smooth Runge function, Lasso’s automatic variable selection provides more robust regularization than Ridge’s uniform shrinkage.

F. Stochastic Gradient Descent Performance Analysis

Table IV compares full-batch and mini-batch gradient descent for OLS regression ($N = 1,000,000$). All methods converge to similar MSE values (0.0268). Ridge and Lasso exhibit comparable patterns with appropriately adjusted learning rates.

Mini-batch SGD converges in dramatically fewer epochs (7-38 vs. 71-163 iterations). With batch size 256 and 750,000 training samples, each epoch performs 2,930 gradient updates versus full-batch’s single update. Adaptive methods (AdaGrad, RMSprop, Adam) demonstrate their designed advantage in stochastic settings with robustly handling noisy gradient estimates without extensive learning rate tuning, though Momentum remains fastest overall.

Despite fewer epochs, mini-batch methods don’t consistently achieve faster wall-clock times. Momentum shows similar timing (0.25s vs. 0.27s), but RMSprop is significantly slower (1.59s vs. 0.38s) due to computational overhead. Data shuffling, batch extraction, and per-parameter learning rate updates across thousands of mini-batches takes up most of the time. For this moderately-sized dataset, overhead negates per-epoch gains.

For convex problems with manageable data sizes, full-batch gradient descent with momentum provides optimal performance. Mini-batch SGD becomes essential for memory-constrained scenarios or non-convex problems, where adaptive methods provide additional robustness to noisy gradients.

G. Bootstrapping

Figure 11 illustrates the Bias-Variance tradeoff for an OLS model with 200 data samples. Using the bootstrap resampling method with 2000 resamples we calculated the averaged Bias², variance and MSE across the 2000 resamples. From the figure we see that as the complexity of the model increase (polynomial degree increase) the bias of the model decreases while the variance increases. Specifically we see that for polynomial degrees in the range $p \in (1, 7)$ the system is dominated by Bias², then there is a stable zone for both bias and variance $p \in (7, 15)$, before variance begins to increase rapidly in the region $p \in (15, 20)$. In ML-terms, this translates to: our model is underfitting in $p \in (1, 7)$ and overfitting in $p \in (15, 20)$. Our model is too simple to capture the underlying pattern (Runge's function) in $p \in (1, 7)$, and our model is too flexible in $p \in (15, 20)$, so it starts to capture the random noise in the data.

Furthermore, we see that the measured MSE curve closely follows the curve of $\text{Var}[\hat{y}] + \text{Bias}^2[\hat{y}]$, and the minimum of this curve is approximately the irreducible error. This is expected. In the theory in Appendix A 1, we showed that the expected error could be decomposed as $E[MSE] = \text{Var}[\hat{y}] + \text{Bias}^2[\hat{y}]$. We also found that the minimum MSE the system could obtain is the irreducible error σ^2 . Our system has noise $\sigma = 0.1$, so the irreducible error $\sigma^2 = 0.01$. Our results seem to agree with theory.

We expanded the parameter space to perform the same bootstrap-analysis for polynomial degrees $p \in (1, 20)$ and sample-sizes $n \in (200, 575)$. In Figure 12 the heatmaps show how MSE, Bias² and variance varies across this parameter space. The results reveal that the best OLS model complexity, across the entire space is at $n = 300$ with degree = 14 (MSE ≈ 0.008). This is the sweet-spot where bias² (0.0076) and variance (0.0008) added together give the lowest error. Since bias dominates this point, we can determine that the model is still slightly underfitting even at degree 14. Now we notice something unexpected: the lowest MSE we found is lower than the irreducible error. The reason we find a result slightly lower than this limit is probably an artifact of the random train-test-splits in the trains-test split. In our bootstrap method we left 25% of the data out for evaluation of the sampled models. This gives $300 \cdot 0.25 = 75$ for testing. This random 75 point subset of the data might not accurately represent the statistical properties of the full Gaussian noise-distribution $\epsilon \sim N(0, \sigma^2)$.

In the figure we also notice another interesting pattern. We see that the MSE, bias and variance follow the same pattern as in Figure 11: bias dominates in regions with low polynomial degree, and variance dominates in regions with high polynomial degree. The third dimension

in the parameter space (sample-sizes) only shifts these regions. Lower sample-sizes makes variance dominate at lower polynomial degrees, and large sample-sizes pushes this to higher polynomial degrees. The Bias seems to be quite unaffected by the sample-sizes. This behavior is expected. Variance is an effect of noise: more data averages out this effect. Bias also slightly decrease with larger sample-sizes. The effect is probably less visible, since Bias measures how well the model can capture the underlying pattern (Runge's function). This is determined by both the model parameters (which is better estimated with more data), but most importantly, by the chosen basis for the model itself. Therefore the bias does not change that much with different sample-sizes. In Figure 13 we visually see the statistical properties we have just discussed. It shows how choosing a too simple model $p = 2$ fundamentally limits how well the model can fit the data. It also illustrates how the variance in the models performance is reduced by choosing a higher sample size $n = 575$.

H. Cross Validation

Figure 14 shows a comparison between bootstrap and 5-fold CV as generalization error estimators for OLS. In bootstrap we split the data in 0.25/0.75 test/train samples. Then we resampled 2000 times to train our OLS model and evaluate it on the fixed 0.25 test data. The plot shows the averaged MSE over the sampled data and predicts how well our model would perform on unseen data. As we discussed in the Bootstrapping section, we see that the model is predicted to under-fit in the regions $p \in (1, 7)$ and overfit in the region $p \in (15, 20)$. We then used 5-fold Cross-Validation as an estimator, where we trained our models on the full data-set and reported the average MSE with standard-deviation across the 5-folds. In the Figure we see that CV-estimator follows the Bootstrap-prediction closely (largely within $\pm 1std$) in the region $p \in (1, 15)$. Then, variance increases in the Bootstrapping estimate from $p = 15$, and the two estimators start to diverge. This is expected. In our implementation of Cross-Validation the model is trained on the full dataset, and then within 5-folds a test-point is chosen on rotation of each fold. In our respectively dense grid $x \in (-1, 1)$ with $n \in (200, 575)$ sample points, each validation point is destined to have near neighbours within the fold. Therefore our implementation of Cross-Validation method might be prone to being too optimistic, and report lower variance estimations. This might be the reason for why the two estimators start to diverge at higher degrees where variance kicks in.

We then wanted to analyse the effectiveness of regularization models (Ridge and Lasso) in the full parameter space we used for bootstrapping, but extended for thoroughness. We performed 5-fold cross-validation for OLS, Ridge (best λ), and Lasso (best λ) across three

noise levels ($\sigma \in \{0.1, 0.2, 0.3\}$). From the previous discussion we knew that models performance could be affected by the random train-test split. Therefore we averaged the MSE across 10 random seeds. For 20 lambdas in the range $\lambda \in (-5, 2)$ in logspace. After around 2 million models and more than one hour of compute we selected the best λ 's independently for each (n, degree) . We then computed the relative difference in percent between OLS and Ridge, and OLS and Lasso respectively in this parameter space. Figure 15 quantifies these improvements: Ridge shows $> 1\%$ improvement in 4%, 13%, and 23% of configurations at $\sigma \in \{0.1, 0.2, 0.3\}$ respectively, with maximum gains of 15-23%. Lasso exhibits similar patterns with maximum gains of 19-22%. We see that the improvement are concentrated at high complexity and high noise, confirming regularization's primary benefit is controlling overfitting when signal-to-noise ratio is low. From the figure we see that the regularization is most effective at higher degrees, in the same region where we observed variance kick-in in the Bootstrap analysis. Furthermore we notice that Ridge generally performs better than Lasso, especially with low noise. This is probably because Lasso over-regulate by completely pushing the coefficients to zero, whereas Ridge only reduce them.

V. CONCLUSION

This study systematically examined Ordinary Least Squares (OLS), Ridge, and Lasso regression for approximating Runge's function under noisy conditions. As shown in our OLS analysis, the training MSE decreases monotonically with model complexity, while the test MSE follows a U-shaped curve, reflecting the classic bias-variance trade-off. Overfitting becomes evident beyond polynomial degree ten, where test performance degrades despite continued improvement on the training data. Introducing ℓ_2 regularization through Ridge regression stabilizes the model by shrinking coefficients, reducing oscillations near the boundaries, and mitigating Runge's phenomenon. Ridge remains stable up to approximately degree 30, offering a smoother and more generalizable fit than OLS. Lasso regression provides even greater robustness, maintaining low test MSE across all examined degrees by setting irrelevant coefficients exactly to zero. Its feature elimination capability prevents high-degree overfitting, making Lasso the most stable of the three. Our gradient descent analysis confirmed that, with properly tuned learning rates, plain gradient descent reproduces analytical OLS and Ridge solutions to numerical precision. Ridge consistently required slightly more iterations due to its regularization

term but benefited from improved numerical conditioning. Among adaptive optimizers, Momentum achieved the fastest convergence (70 iterations) and Adam demonstrated the most reliable stability across learning rates. For these smooth, deterministic cost functions, adaptive methods offered only modest performance improvements, serving mainly to enhance robustness to hyperparameter selection. Stochastic gradient descent reached comparable MSE values with dramatically fewer epochs, but computational overhead offset any runtime gains for our dataset sizes. Resampling methods reinforced these findings. Bootstrap analysis identified the optimal OLS model complexity around polynomial degree 14 with approximately 300 data points, matching the bias-variance decomposition where bias and variance are balanced. Cross-validation results showed that as noise increased, the advantages of regularization became more pronounced: Ridge and Lasso achieved up to 20% lower MSE than OLS at high noise and high complexity. Together, these results confirm that Ridge and Lasso effectively counteract Runge's phenomenon by stabilizing high-degree polynomial fits—Ridge through coefficient shrinkage and Lasso through feature selection. Overall, Lasso regression emerges as the most balanced and robust model, maintaining low variance, good generalization, and strong performance across noise levels and polynomial complexities.

A. Further studies

Our bias-variance analysis reveals distinct performance patterns. At low noise ($\sigma = 0.1$), Lasso excels at around the first three polynomial degrees, while OLS dominates from polynomial degree 10 to 20. At higher noise ($\sigma \geq 0.2$), Lasso consistently outperforms OLS at high polynomial degree at low sample sizes.

This behavior may stem from Runge's symmetry favoring even-power terms (x^2, x^4, \dots), which Lasso's feature selection exploits by retaining relevant coefficients while eliminating noise-fitting terms. Future work could explore asymmetric test functions and elastic net regression, which combines Ridge's shrinkage with Lasso's feature selection.

The fundamental concepts explored here, gradient descent optimization, regularization techniques, and bias-variance trade-offs, form the foundation of modern AI systems. Large language models like GPT and Claude employ these same principles: gradient-based optimization trains billions of parameters, while regularization prevents overfitting. Understanding polynomial regression provides essential intuition for deep learning and contemporary AI applications.

[1] Wikipedia contributors. Runge's phenomenon. https://en.wikipedia.org/wiki/Runge%27s_phenomenon, 2025.

Accessed: 2025-10-05.

- [2] H. Haug, S. S. Thommesen, C. A. Falchenberg, and L. L. Storborg. Project 1. <https://github.com/livelstorborg/FYS-STK4155>, 2025. GitHub repository.
- [3] J. D. Hunter. Matplotlib: A 2d graphics environment. *Computing in Science & Engineering*, 9(3):90–95, 2007.
- [4] Charles R. Harris, K. Jarrod Millman, Stéfan J. van der Walt, Ralf Gommers, Pauli Virtanen, David Cournapeau, Eric Wieser, Julian Taylor, Sebastian Berg, Nathaniel J. Smith, Robert Kern, Matti Picus, Stephan Hoyer, Marten H. van Kerkwijk, Matthew Brett, Allan Haldane, Jaime Fernández del Río, Mark Wiebe, Pearu Peterson, Pierre Gérard-Marchant, Kevin Sheppard, Tyler Reddy, Warren Weckesser, Hameer Abbasi, Christoph Gohlke, and Travis E. Oliphant. Array programming with NumPy. *Nature*, 585(7825):357–362, September 2020. <https://doi.org/10.1038/s41586-020-2649-2>.
- [5] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.
- [6] Michael L. Waskom. seaborn: statistical data visualization. *Journal of Open Source Software*, 6(60):3021, 2021. <https://doi.org/10.21105/joss.03021>.
- [7] The pandas development team. pandas-dev/pandas: Pandas, August 2025. <https://doi.org/10.5281/zenodo.16918803>.

Appendix A

1. Bias-Variance Decomposition

The cost function $C(X, \theta)$ can be expressed in terms of the variance-bias decomposition. This can be derived as follows.

We start by defining Variance and Bias².

$$\text{Var}[\tilde{y}(x)] = \mathbb{E}[(\tilde{y}(x) - \mathbb{E}[\tilde{y}(x)])^2], \quad (\text{A1})$$

$$\text{Bias}_f^2[\tilde{y}(x)] = (\mathbb{E}[\tilde{y}(x)] - f(x))^2. \quad (\text{A2})$$

We assume that the expectation values are taken over different samples for the training data. For fixed values x , we look at how the prediction \tilde{y} changes when the model is trained on different training samples (e.g. bootstrap). Rewriting the terms inside the parenthesis.

$$\begin{aligned} (y - \tilde{y})^2 &= (y - \tilde{y} + E[\tilde{y}] - E[\tilde{y}])^2 \\ &= (y - E[\tilde{y}])^2 + (\tilde{y} - E[\tilde{y}])^2 \\ &\quad - 2(y\tilde{y} - yE[\tilde{y}] - \tilde{y}E[\tilde{y}] + E[\tilde{y}]^2) \end{aligned}$$

Now distributing the expectation operator linearly:

$$\begin{aligned} E[(y - \tilde{y})^2] &= E[(y - E[\tilde{y}])^2] + E[(\tilde{y} - E[\tilde{y}])^2] \\ &\quad - 2(E[y\tilde{y}] - E[y]E[\tilde{y}] - E[\tilde{y}]E[\tilde{y}] + E[\tilde{y}]^2) \\ &= E[(y - E[\tilde{y}])^2] + E[(\tilde{y} - E[\tilde{y}])^2] \\ &\quad - 2(E[y\tilde{y}] - E[y]E[\tilde{y}]) \end{aligned}$$

Now we assume that our sampled values y , can be expressed as $y = f(x) + \epsilon$, where $f(x)$ is some deterministic function and $\epsilon \sim N(0, \sigma^2)$ is assumed to be a Gaussian distribution of random noise that is independent of $f(x)$. The noise in the training data is independent of the noise in the test samples. Since ϵ represents the noise in y and is independent of the training-data used to construct \tilde{y} , ϵ is independent of \tilde{y} : $E[\epsilon\tilde{y}] = E[\epsilon]E[\tilde{y}]$. ϵ follows a normal distribution with zero mean and variance σ^2 , so per definition $E[\epsilon] = 0$, and $\text{Var}(\epsilon) = E[\epsilon^2] = \sigma^2$. We also notice that since $f(x)$ is deterministic, we know $E[f(x)] = f(x)$ and that $E[f(x)\tilde{y}] = f(x)E[\tilde{y}]$. We can now start by expanding the expression within the parenthesis and use the substitution $y = f(x) + \epsilon$ to simplify.

$$\begin{aligned} (E[y\tilde{y}] - E[y]E[\tilde{y}]) &= E[(f + \epsilon)\tilde{y}] - E[f + \epsilon]E[\tilde{y}] \\ &= E[f]E[\tilde{y}] + E[\epsilon]E[\tilde{y}] - E[f]E[\tilde{y}] \\ &\quad - E[\epsilon]E[\tilde{y}] \\ &= 0 \end{aligned}$$

So we are left with:

$$\begin{aligned} E[(y - \tilde{y})^2] &= E[(y - E[\tilde{y}])^2] + E[(\tilde{y} - E[\tilde{y}])^2] \\ &\quad - 2(E[y\tilde{y}] - E[y]E[\tilde{y}]) \\ &= E[(y - E[\tilde{y}])^2] + E[(\tilde{y} - E[\tilde{y}])^2] \end{aligned}$$

Here we notice that this is the sum of the variance and the bias squared, where the bias squared is given by:

$$\text{Bias}_y^2[\tilde{y}] = E[(y(x) - E[\tilde{y}])^2]$$

$y(x) = f(x) + \epsilon$ again.

$$\begin{aligned} E[(y - \tilde{y})^2] &= \text{Var}[\tilde{y}] + \text{Bias}_y^2[\tilde{y}] \\ &= E[(y - E[\tilde{y}])^2] + E[(\tilde{y} - E[\tilde{y}])^2] \\ &= E[(\tilde{y} - E[\tilde{y}])^2] + E[y^2] - 2E[yE[\tilde{y}]] + E[\tilde{y}]^2 \\ &= E[(\tilde{y} - E[\tilde{y}])^2] + E[f^2] + E[2f\epsilon] + E[\epsilon^2] \\ &\quad - 2(E[f]E[\tilde{y}] + E[\epsilon]E[\tilde{y}]) + E[\tilde{y}]^2 \\ &= E[(\tilde{y} - E[\tilde{y}])^2] + E[f^2] - 2E[f]E[\tilde{y}] \\ &\quad + E[\tilde{y}]^2 + E[\epsilon^2] \\ &= E[(\tilde{y} - E[\tilde{y}])^2] + E[(f - E[\tilde{y}])^2] + E[\epsilon^2] \\ &= \text{Var}[\tilde{y}] + \text{Bias}_f^2[\tilde{y}] + \sigma^2 \end{aligned}$$

Therefore we have shown that the cost-function or the expected error, can be decomposed as:

$$\begin{aligned} C(X, \theta) &= E[(y - \tilde{y})^2] = \text{Var}[\tilde{y}] + \text{Bias}_y^2[\tilde{y}] \\ &= \text{Var}[\tilde{y}] + \text{Bias}_f^2[\tilde{y}] + \sigma^2 \end{aligned}$$

We see that we can express the bias-term both using the true function $f(x)$ and the values we sample y . These values differ by the irreducible error σ^2 .

We can expand this as before using the substitution